

# High Performance LLMs From First Principles (2024)

## This session:

- Implement efficient inference!
  - What else to learn? (lots!)
- One deep dive: talk numerics

## Next/last session:

- Pallas deep dive with Sharad Vikram!

# My Asks

Please ask lots of questions! Just raise your hand or speak up!

If there are topics you're interested in, message me between sessions.

Join the discord! <https://discord.gg/2AWcVatVAw>

Do the exercises! Give feedback, ask questions!

Website: <https://github.com/rwitten/HighPerfLLMs2024>

# High Performance Inference (Generate, Batch=1)

- Pass in single token, not multiple tokens.
  - Deal with embedding lookup (`jax.lax.dynamic_slice_in_dim`)
  - Now NN is silly.
- KV Cache
  - Create KV Cache
  - Read/write from KV cache (`jax.lax.dynamic_update_index_in_dim`)
  - Fix Attention to work with KV cache
- Time it and compare to memory bandwidth limit.
- Todo (not in class):
  - Quantization!
  - Prefill
  - Batch > 1 (as we saw on rooflines, batch > 1 can give more throughput at the same latency)


# Much Much More To Learn!

- Thinking about Performance goes much deeper!
  - Next week, Sharad Vikram, co-inventor of the kernel language Pallas, will explain Pallas!
  - Pallas lets you more directly program the GPUs and TPUs to achieve high performance in cases where pure XLA can't.
- Broadly the other category is ML Modeling (which we haven't covered at all)!

Many many questions!

- How should we design the network? What types of layers, what order, what mix and what sizes?
- How should we quantize it?
- What data should we use?
- What optimizer should we use?

# ML Modeling – try things, measure loss



- Typical graph! “Our method” looks good!

[Thanks Xu et Al, 2017 for the randomly selected example graph.](#)

# ML Modeling – minimizing loss is useful!

TEXT			Gemini Ultra	GPT-4
Capability	Benchmark Higher is better	Description		API numbers calculated where reported numbers were missing
General	MMLU	Representation of questions in 57 subjects (incl. STEM, humanities, and others)	90.0% CoT@32*	86.4% 5-shot** (reported)
Reasoning	Big-Bench Hard	Diverse set of challenging tasks requiring multi-step reasoning	83.6% 3-shot	83.1% 3-shot (API)
	DROP	Reading comprehension (F1 Score)	82.4 Variable shots	80.9 3-shot (reported)
	HellaSwag	Commonsense reasoning for everyday tasks	87.8% 10-shot*	95.3% 10-shot* (reported)
Math	GSM8K	Basic arithmetic manipulations (incl. Grade School math problems)	94.4% maj1@32	92.0% 5-shot CoT (reported)

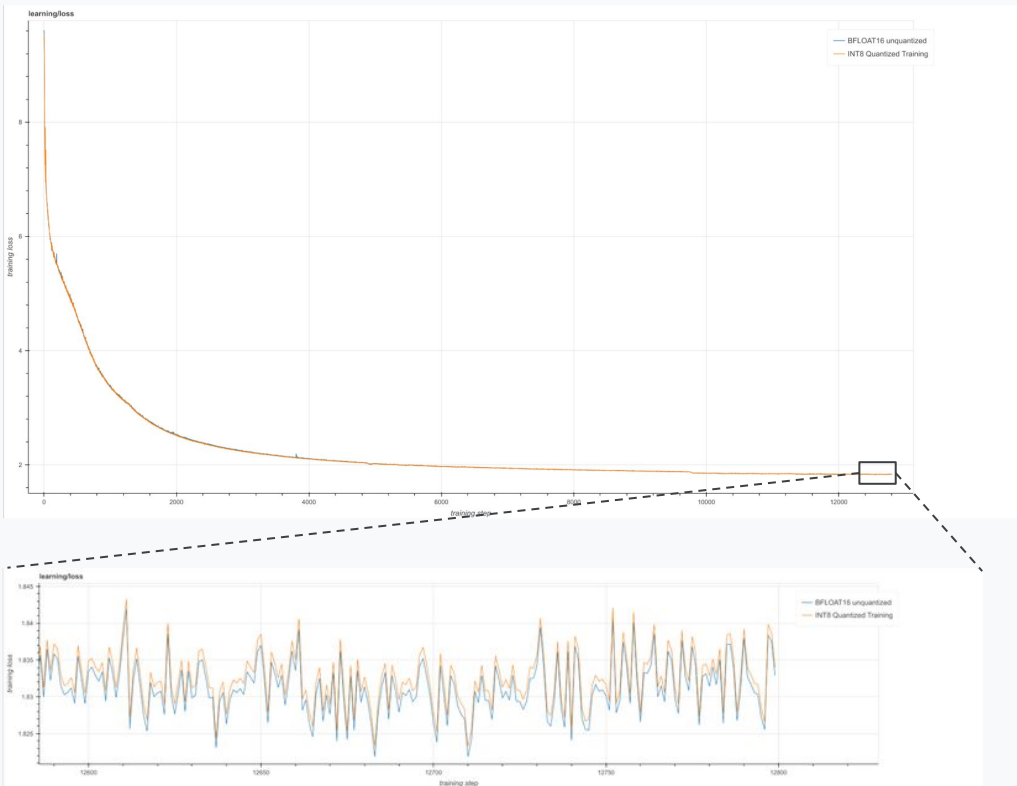
# Deep dive into “Numerics” – typical mix of Performance and Modeling!

- Generally, it is much faster for computer chips to do math on lower bit representations of numbers than on higher bit representations!
- But lower precision introduces errors so the models learn less quickly (or maybe not at all).
- The first big change was going from float32 (32-bit floats) to bfloat16 (16-bit floats).
  - (Introduced in 2019 by Google!)
- As an example, let's consider summing up a vector.
  - Coding exercise...



# Numerics go deep!

- Nowadays, we're training and inferring on models with 8-bit matrix multiplies!
  - For training a ~1.5x speedup, very small quality gap!
- Typically int8 on TPU and fp8 on GPU.
- ([Source Link](#))



**Next/last session:**

- **Pallas deep dive with Sharad Vikram!**

**Thanks!**

**Ping me ([rwitten@google.com](mailto:rwitten@google.com)) with  
feedback, suggested topics, etc!**